

TRANSACTION CONCEPT

In database management systems (DBMS), the concept of a transaction plays a crucial role in ensuring data integrity, consistency, and reliability. Here's a detailed explanation of what a transaction is and its importance:

1. **Definition:** A transaction in DBMS refers to a logical unit of work that is performed within a database management environment. It represents a sequence of operations that are treated as a single unit.
2. **Properties of Transactions:**
 - **Atomicity:** Transactions are atomic, meaning they are either completed in full or not executed at all. If any part of a transaction fails (due to system error, user cancellation, or other reasons), the entire transaction is rolled back to its original state (the state it was in before the transaction began).
 - **Consistency:** Transactions ensure that the database remains in a consistent state before and after the transaction. This means that the integrity constraints and business rules are enforced during the execution of the transaction.
 - **Isolation:** Transactions operate independently of each other. Changes made by one transaction are typically not visible to other transactions until they are committed. This ensures that transactions do not interfere with each other and helps maintain data integrity.
 - **Durability:** Once a transaction is committed (its changes are successfully written to the database), the changes persist even in the event of a system failure. This is achieved through mechanisms like transaction logging and database recovery techniques.

Importance of Transactions

1. **Data Integrity:** Transactions ensure that database operations maintain the correctness and consistency of the data. The atomicity property prevents incomplete or partial updates that could leave the database in an inconsistent state.

2. **Concurrency Control:** Transactions facilitate concurrent access to the database by multiple users or applications while ensuring that their operations do not interfere with each other. Isolation mechanisms prevent dirty reads, non-repeatable reads, and other anomalies that can occur due to concurrent access.
3. **Recovery and Undo Operations:** Transactions provide a mechanism for recovery in case of failures. The logging of transaction operations allows DBMS to roll back or undo changes made by incomplete or failed transactions, ensuring data durability.
4. **Application Development:** Transactions simplify application development by allowing developers to group database operations into logical units. This simplification makes it easier to manage complex interactions with the database and maintain data integrity.

EXAMPLES

Example: Online Banking Transfer

Alice initiates a \$100 transfer from her checking account to Bob's savings account:

- System checks Alice's account balance.
- System deducts \$100 from Alice's account and adds it to Bob's account.
- Transaction is committed if successful; otherwise, it's rolled back to maintain data integrity.

This transaction ensures that the transfer is completed atomically (all-or-nothing), maintains consistency of account balances, isolates changes from other transactions, and ensures durability of the transaction once committed.

Example: Online Order Processing

Alice places an online order for a laptop:

- System checks product availability and deducts the ordered quantity from inventory.
- System updates Alice's order history and deducts payment from her account.

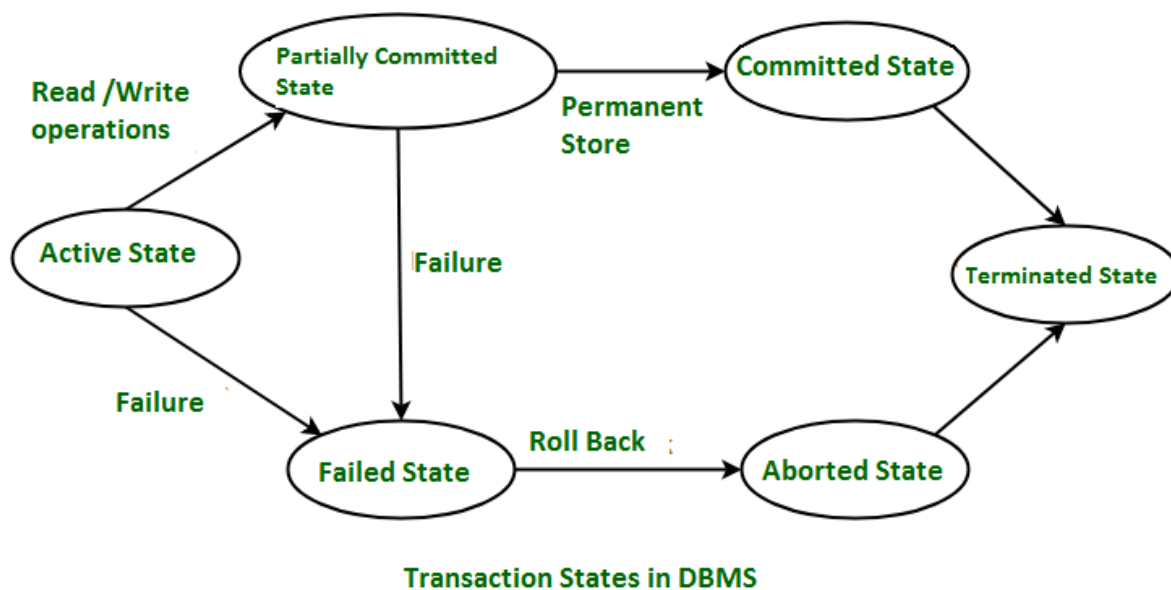
- Transaction is committed if all steps succeed; otherwise, it's rolled back to maintain consistency.

This transaction ensures that the order processing is completed atomically, maintains inventory accuracy, updates customer records, and ensures financial transactions are accurate and complete.

TRANSACTION STATE

States through which a transaction goes during its lifetime. These are the states which tell about the current state of the Transaction and also tell how we will further do the processing in the transactions. These states govern the rules which decide the fate of the transaction whether it will commit or abort.

They also use Transaction log. Transaction log is a file maintain by recovery management component to record all the activities of the transaction. After commit is done transaction log file is removed.



These are different types of Transaction States:

1. **Active State –**

When the instructions of the transaction are running then the transaction is in active state. If all the 'read and write' operations are performed without any error then it goes to the "partially committed state"; if any instruction fails, it goes to the "failed state".

2. **Partially Committed –**

After completion of all the read and write operation the changes are made in main memory or local buffer. If the changes are made permanent on the Database then the state will change to "committed

state” and in case of failure it will go to the “failed state”.

3. Failed State –

When any instruction of the transaction fails, it goes to the “failed state” or if failure occurs in making a permanent change of data on Data Base.

4. Aborted State –

After having any type of failure the transaction goes from “failed state” to “aborted state” and since in previous states, the changes are only made to local buffer or main memory and hence these changes are deleted or rolled-back.

5. Committed State –

It is the state when the changes are made permanent on the Data Base and the transaction is complete and therefore terminated in the “terminated state”.

6. Terminated State –

If there isn't any roll-back or the transaction comes from the “committed state”, then the system is consistent and ready for new transaction and the old transaction is terminated.

LOCK-BASED PROTOCOLS

Lock-based protocols in database management systems (DBMS) are mechanisms used to manage concurrent access to shared data to ensure data consistency and prevent conflicts between transactions. Here's a detailed explanation of lock-based protocols:

Purpose of Lock-Based Protocols

Concurrency control in DBMS aims to allow multiple transactions to access and modify data concurrently while ensuring that the final state of the database remains consistent. Lock-based protocols achieve this by controlling access to shared resources (data items) using locks.

Types of Locks

1. Shared (Read) Lock (S):

- Allows multiple transactions to read a data item simultaneously.
- Transactions holding a shared lock can only read the data item but cannot modify it.

2. Exclusive (Write) Lock (X):

- Allows a transaction to exclusively write (modify) a data item.
- Only one transaction can hold an exclusive lock on a data item at a time.

Lock-Based Protocols

1. Two-Phase Locking (2PL):

- **Growing Phase:** Transactions acquire locks on needed data items but cannot release any locks.
- **Shrinking Phase:** Transactions release locks but cannot acquire any new locks.
- Guarantees serializability (equivalent to executing transactions sequentially) and prevents conflicts (e.g., lost updates, inconsistent reads).

2. Strict Two-Phase Locking (Strict 2PL):

- A variant of 2PL where transactions hold all their exclusive locks until they commit or abort.
- Ensures that no other transaction can access the locked data until the lock holder releases the locks, thereby preventing cascading aborts.

3. Deadlock Handling:

- DBMS must handle deadlock situations where transactions wait indefinitely for locks held by each other.
- Techniques include deadlock detection and resolution strategies (e.g., timeout mechanisms, deadlock detection algorithms, and deadlock prevention through lock ordering).

Example Scenario

Consider two transactions in an inventory management system:

- **Transaction T1:** Updates the quantity of a product in the inventory.
- **Transaction T2:** Calculates the total value of all products in the inventory.

Execution with Lock-Based Protocols:

- T1 requests an exclusive lock (X lock) on the inventory data item it wants to update.
- T2 requests a shared lock (S lock) on the inventory data item it wants to read.
- T1 completes its update and releases the exclusive lock.
- T2 completes its calculation and releases the shared lock.

Advantages of Lock-Based Protocols

- **Concurrency Control:** Ensures that transactions execute in a controlled manner to prevent conflicts and maintain data consistency.
- **Isolation:** Allows transactions to execute independently without interfering with each other's operations.

- **Performance:** Efficiently manages concurrent access to data, improving system throughput and response time.

Considerations and Limitations

- **Deadlocks:** Requires mechanisms to detect and resolve deadlocks to avoid system hangs.
- **Overhead:** Managing locks incurs overhead in terms of processing and memory resources.
- **Complexity:** Designing and implementing effective lock-based protocols requires careful consideration of transaction behavior and system performance.

In summary, lock-based protocols are essential for ensuring concurrent transactions in DBMS maintain data integrity and consistency. They provide mechanisms to control access to shared resources through locks, enabling efficient and reliable transaction processing in multi-user environments.

VALIDATION - BASED PROTOCOLS

Validation Based Protocol is also called Optimistic Concurrency Control Technique. This protocol is used in DBMS (Database Management System) for avoiding concurrency in transactions. It is called optimistic because of the assumption it makes, i.e. very less interference occurs, therefore, there is no need for checking while the transaction is executed.

In this technique, no checking is done while the transaction is been executed. Until the transaction end is reached updates in the transaction are not applied directly to the database. All updates are applied to local copies of data items kept for the transaction. At the end of transaction execution, while execution of the transaction, a **validation phase** checks whether any of transaction updates violate serializability. If there is no violation of serializability the transaction is committed and the database is updated; or else, the transaction is updated and then restarted.

Key Concepts and Steps Involved

1. Execution Phase:

- Transactions execute without acquiring locks on data items. They read and write data as necessary.

2. Validation Phase:

- After a transaction has completed its execution phase and is ready to commit, it undergoes a validation step.
- During validation, the DBMS checks if the transaction's updates conflict with updates made by other concurrent transactions that have committed since the transaction began.

3. Validation Criteria:

- **Conflict Detection:** The DBMS checks for conflicts between the current transaction and other committed transactions. Typical conflicts include:
 - **Read-Write Conflict:** If another transaction has written to a data item that the current transaction has read.

- **Write-Write Conflict:** If another transaction has written to a data item that the current transaction also intends to write.
- If no conflicts are detected during validation, the transaction is considered valid and can proceed to commit.

4. Commit Phase:

- Validated transactions proceed to commit their changes to the database.
- Once committed, the changes become permanent and durable.

5. Abort and Restart:

- If conflicts are detected during validation, the transaction is aborted.
- The aborted transaction may need to be restarted or rolled back to ensure data consistency. In order to perform the Validation test, each transaction should go through the various phases as described above. Then, we must know about the following three time-stamps that we assigned to transaction T_i , to check its validity:
 - **1. Start(T_i):** It is the time when T_i started its execution.
 - **2. Validation(T_i):** It is the time when T_i just finished its read phase and begin its validation phase.
 - **3. Finish(T_i):** the time when T_i end it's all writing operations in the database under write-phase.
- Two more terms that we need to know are:
 - **1. Write_set:** of a transaction contains all the write operations that T_i performs.
 - **2. Read_set:** of a transaction contains all the read operations that T_i performs.

Example Scenario

Consider two transactions in an inventory management system:

- **Transaction T1:** Reads the quantity of a product and then updates its price.
- **Transaction T2:** Reads the updated price of the product and then calculates its total value.

Execution with Validation-Based Protocol:

- T1 completes its execution and proposes its update (price change).
- T2 begins execution after T1 and reads the updated price.
- During T2's validation phase, the DBMS checks if the update proposed by T1 conflicts with T2's read.
- If T2's read of the updated price is consistent (no conflict), T2 proceeds to commit. Otherwise, T2 aborts.

Advantages of Validation-Based Protocols

- **Reduced Lock Contention:** Compared to lock-based protocols, validation-based protocols can reduce contention on locks, leading to better concurrency and system throughput.
- **Deadlock Avoidance:** Since transactions do not acquire locks, deadlock scenarios are less likely to occur.
- **Efficiency:** Allows transactions to execute without blocking each other, improving system performance in environments with high concurrency.

Considerations and Limitations

- **Complexity:** Implementing validation-based protocols requires sophisticated conflict detection mechanisms and careful transaction scheduling.
- **Overhead:** The validation phase may introduce overhead, especially in scenarios with high transaction rates or complex data dependencies.
- **Conflict Resolution:** Efficient conflict detection and resolution strategies are crucial for maintaining data consistency and avoiding anomalies.

In summary, validation-based protocols offer an alternative approach to concurrency control in DBMS, emphasizing conflict detection during a transaction's validation phase rather than relying on locks. They aim to achieve high concurrency and performance while ensuring transaction serializability and data integrity.

TIMESTAMP - BASED PROTOCOLS

Timestamp-based protocols are concurrency control mechanisms used in database management systems (DBMS) to ensure serializability of transactions based on their timestamps. These protocols use timestamps (usually based on transaction start times or commit times) to determine the order of transactions and enforce isolation levels. Here's a detailed explanation of timestamp-based protocols:

Key Concepts of Timestamp-Based Protocols

1. Timestamp Assignment:

- Each transaction is assigned a unique timestamp when it begins execution or when it is first submitted to the system.
- Timestamps can be assigned based on system clocks or using a monotonically increasing counter.

2. Transaction States:

- **Active:** The transaction is currently executing.
- **Blocked:** The transaction is waiting for a lock or resource.
- **Committed:** The transaction has successfully completed its execution and committed its changes to the database.
- **Aborted:** The transaction encountered an error or was rolled back due to conflicts.

3. Concurrency Control Mechanisms:

- **Timestamp Ordering:** Transactions are ordered based on their timestamps to enforce a serializable schedule.
- **Transaction Conflict Resolution:** Conflicts between transactions are resolved based on their timestamps:
 - **Read-Write Conflict:** A transaction with an older timestamp is allowed to read a data item that a transaction with a newer timestamp has written.

- **Write-Write Conflict:** A transaction with an older timestamp is allowed to write to a data item that a transaction with a newer timestamp has also written.
- Conflicts where timestamps are equal can be resolved using additional criteria such as transaction ID or using a tie-breaking mechanism.

4. Transaction Execution:

- **Read and Write Operations:** Transactions perform read and write operations on data items without acquiring locks.
- **Timestamp Comparison:** Before performing a read or write operation on a data item, a transaction compares its timestamp with the timestamp of the last transaction that accessed or modified the data item.

5. Transaction Commit:

- When a transaction completes its operations, it attempts to commit its changes.
- The DBMS checks if the transaction's timestamp is still valid (has not been surpassed by a newer transaction) before allowing it to commit.

6. Transaction Rollback:

- If a transaction encounters a conflict (e.g., read-write or write-write conflict), it may be aborted and rolled back to maintain consistency.
- Aborted transactions release any locks they hold and undo their changes to the database.

Example Scenario

Consider two transactions in an online shopping system:

- **Transaction T1:** Places an order for a product and updates the inventory.
- **Transaction T2:** Updates the price of the product.

Execution with Timestamp-Based Protocol:

- T1 begins execution and is assigned timestamp $T1_{start}$.
- T2 begins execution and is assigned timestamp $T2_{start}$.
- During execution, T2 attempts to update the product price.
- The DBMS checks if $T2_{start}$ is greater than $T1_{start}$ to determine if T2 can proceed (assuming no other conflicts).

Advantages of Timestamp-Based Protocols

- **High Concurrency:** Transactions can execute concurrently without blocking each other, leading to improved system performance.
- **Deadlock Prevention:** Since transactions do not acquire locks, deadlock situations are inherently avoided.
- **Efficiency:** Lower overhead compared to lock-based protocols because transactions do not need to manage locks.

Considerations and Limitations

- **Clock Synchronization:** Requires synchronized clocks across all nodes in distributed systems to ensure correct timestamp ordering.
- **Timestamp Granularity:** Determining an appropriate timestamp granularity (e.g., millisecond, microsecond) is crucial to avoid conflicts.
- **Correctness:** Careful implementation is required to ensure that transactions are correctly ordered based on their timestamps to maintain data consistency.

In summary, timestamp-based protocols provide an efficient mechanism for concurrency control in DBMS by using transaction timestamps to order and schedule transactions. They facilitate high concurrency while ensuring serializability and data consistency, making them suitable for environments with a large number of concurrent transactions.

RECOVERY AND ATOMICITY

In database management systems (DBMS), recovery and atomicity are closely intertwined concepts that ensure the durability and consistency of data despite system failures or crashes. Let's explore how recovery mechanisms uphold atomicity and ensure data integrity:

Atomicity and its Role in Recovery

Atomicity in DBMS refers to the all-or-nothing property of transactions. It ensures that either all operations of a transaction are successfully completed and applied to the database, or none of them are. This property is crucial for maintaining data consistency and integrity, especially during system failures.

1. **Transaction Logs:** DBMS maintains a transaction log, which records all changes made by transactions before they are applied to the actual database. This log serves as a record of operations that must be undone (rolled back) or redone (reapplied) during recovery.
2. **Write-Ahead Logging (WAL):** This technique ensures that changes made by transactions are first recorded in the transaction log before they are applied to the database. This guarantees that if a transaction commits, its changes are durable and can be recovered even if a crash occurs before the changes are written to the database.

Recovery Mechanisms

1. **Rollback (Undo) Recovery:**
 - If a transaction fails or encounters an error, all its changes that were recorded in the transaction log but not yet applied to the database are undone.
 - The DBMS uses the undo information in the transaction log to rollback the changes made by the incomplete or failed transaction, ensuring atomicity.
2. **Redo (Forward) Recovery:**
 - After a crash, the DBMS uses the transaction log to redo (reapply) changes made by transactions that were committed but not yet written to the database.

- This process ensures that committed transactions, which may have had their changes lost due to the crash, are re-executed and their effects are reapplied to maintain atomicity and durability.

3. Transaction Commit Protocol:

- Before a transaction is committed, the DBMS ensures that all changes made by the transaction (including metadata in the transaction log) are successfully written to non-volatile storage (disk).
- This step ensures that once a transaction commits, its changes are durable and can survive system failures.

Example Scenario

Consider a banking application where a customer transfers money between accounts:

- **Transaction T1:** Deducts \$100 from Alice's account.
- **Transaction T2:** Credits \$100 to Bob's account.

If a crash occurs after T1 commits but before T2 commits:

- During recovery, the DBMS uses the transaction log to redo the changes made by T2 (credit \$100 to Bob's account) to ensure that the transaction's atomicity is maintained.
- If the crash had occurred before T1 committed, the DBMS would have used the transaction log to undo the changes made by T1 (deduct \$100 from Alice's account) to maintain atomicity.

Advantages and Considerations

- **Data Integrity:** Recovery mechanisms ensure that committed transactions maintain their atomicity, preventing partial updates that could lead to inconsistencies.
- **Performance Impact:** The overhead of maintaining and processing transaction logs for recovery purposes can impact system performance.

- **Complexity:** Implementing robust recovery mechanisms requires careful design and testing to handle various failure scenarios while maintaining data integrity.

In summary, recovery mechanisms in DBMS play a crucial role in ensuring atomicity by using transaction logs to rollback or redo changes made by transactions during system failures or crashes. This ensures that transactions are either fully committed or fully aborted, maintaining the integrity and consistency of the database.

LOG – BASED RECOVERY

Log-based recovery in database management systems (DBMS) is a fundamental technique used to ensure data durability and maintain transaction atomicity and consistency in the event of system failures or crashes. Here's a detailed explanation of how log-based recovery works:


Components of Log-Based Recovery

1. Transaction Log:

- **Purpose:** The transaction log is a sequential record of all actions (operations) performed by transactions that modify the database.
- **Contents:** Each log entry typically includes:
 - **Transaction ID:** Identifies the transaction associated with the log entry.
 - **Operation:** Specifies the type of operation (e.g., read, write) performed by the transaction.
 - **Old Value:** The value of the data item before the operation (for undo purposes).
 - **New Value:** The value of the data item after the operation (for redo purposes).
 - **Commit Record:** Marks the point where a transaction commits, indicating that its changes are durable.

2. Write-Ahead Logging (WAL):

- **Principle:** Changes made by transactions are first recorded in the transaction log before they are applied to the database itself.
- **Ensures Durability:** This ensures that if a transaction commits, its changes are logged and can be recovered even if a system failure occurs before the changes are written to the database.



Log and log records

The log is a sequence of log records, recording all the updated activities in the database. In stable storage, logs for each transaction are maintained. Any operation which is performed on the database is recorded on the log. Prior to performing any modification to the database, an updated log record is created to reflect that modification. An update log record represented as: $\langle T_i, X_j, V_1, V_2 \rangle$ has these fields:

1. **Transaction identifier:** Unique Identifier of the transaction that performed the write operation.
2. **Data item:** Unique identifier of the data item written.
3. **Old value:** Value of data item prior to write.
4. **New value:** Value of data item after write operation.

Other types of log records are:

1. **$\langle T_i \text{ start} \rangle$:** It contains information about when a transaction T_i starts.
2. **$\langle T_i \text{ commit} \rangle$:** It contains information about when a transaction T_i commits.
3. **$\langle T_i \text{ abort} \rangle$:** It contains information about when a transaction T_i aborts.

Steps in Log-Based Recovery

1. **Analysis Phase:**

- **Purpose:** During system recovery (after a crash), the DBMS first scans through the transaction log from the most recent checkpoint to identify transactions that were active (committed but not yet written to the database) at the time of the crash.
- **Identifies Active Transactions:** Transactions marked with commit records but not yet flushed to disk are identified as potentially needing redo operations.

2. Redo Phase:

- **Purpose:** The DBMS re-applies (redoes) changes recorded in the log to ensure that all committed transactions have their effects applied to the database.
- **Process:** It starts from the last checkpoint and applies all logged changes (write operations) of transactions that committed but were not yet written to the database at the time of the crash.
- **Ensures Atomicity and Durability:** Redo phase ensures that committed transactions' changes are durable and not lost due to system failure.

3. Undo Phase:

- **Purpose:** If necessary (e.g., for transactions that were active but not yet committed at the time of the crash), the DBMS uses the transaction log to undo changes that were applied to the database but not yet committed (due to the crash).
- **Rollback Transactions:** Transactions that were active but did not commit are rolled back to ensure atomicity and consistency.
- **Restores Database to a Consistent State:** Undo phase ensures that any partial updates or uncommitted changes are reverted, restoring the database to a consistent state before the crash.

Example Scenario

Consider a scenario where a transaction T1 modifies a record R1, and the DBMS crashes before the change to R1 is written to the database:

- **Redo Phase:** During recovery, the DBMS scans the transaction log, identifies the update operation of T1 on R1, and reapplies this change to R1 to ensure durability.
- **Undo Phase:** If a transaction T2 committed after modifying R2 but before the crash, and T2's changes to R2 were not flushed to disk, the DBMS may undo T2's changes using the transaction log to maintain atomicity and consistency.

Advantages and Considerations

- **Data Durability:** Ensures that all committed transactions' changes are durable and survive system failures.
- **Transaction Atomicity:** Guarantees that transactions are either fully committed or fully rolled back, maintaining data integrity.
- **Performance Overhead:** Maintaining and processing transaction logs for recovery purposes can impose overhead on system performance.
- **Complexity:** Implementing and managing log-based recovery mechanisms requires careful design and testing to handle various failure scenarios effectively.

Log-based recovery is a cornerstone of ensuring reliability and consistency in modern DBMS, providing a robust mechanism to handle failures while preserving data integrity through transaction logging and recovery procedures.